

FROM VOICE NOTE TO FINAL CUT.

Film Cohort, Session 1, distilled to the AI filmmaking guidance.

Saturday, October 3, 2026 · 2.5-hour session led by **Darion D'Anjou** · Case study: the **Modern House** action short

THE SESSION IN SEVEN IDEAS

WHAT TO REMEMBER / WHERE TO START

- 01 Story first, then script**
Dictate the idea, let the LLM outline it and format the screenplay. Keep a real script; camera cues come later.
- 02 Build the world before the shots**
An LLM can model your location in Blender from a reference video. That 3D set drives blocking and continuity.
- 03 Prompt less, then direct**
Models have studied the masters. See what they give you for free, then add control where the story needs it.
- 04 Foundations decide everything**
Style frame, palette and character references carry into every shot. One weak reference can cost days.
- 05 Map every reference**
Number your images and give the model a legend, so no character, prop or location gets swapped.
- 06 Cover the last 20% with inserts**
Generate 30 seconds of 1–2-second cutaways and intercut them to hide janky moments.
- 07 Direct like a studio**
Let AI be your assistant editor and crew. Keep project memory in .md files. Expect tools to change mid-project.

CASE STUDY

MODERN HOUSE

A ~90-second action short:
Darion's Maserati crashes through a glass-and-concrete house, then a one-minute fight.

8 hrs

Claude building the house in Blender on its own, from 100+ frames of a home-tour video

8

thugs designed from one general prompt

19

photoreal keyframes in a neon-noir look

30

reference images Seedance 2.5 accepts in one prompt

2 × 30 s

clips joined with extend-forwards for the fight

480p

test resolution before paying for 1080p

THE CORE LOOP

Script > 3D set > Blocking > Look > Cast > Keyframes > Video > Inserts > Edit

THE WORKFLOW

TWELVE STEPS, IDEA TO CUT.

The pipeline Darion demonstrated on Modern House. The stages of a traditional production stay; the tools change.

PRE-PRODUCTION

STORY / WORLD / LOOK

01 Dictate the story



Record the idea as a voice note. Paste the transcript into ChatGPT, shape an outline, then ask for a properly formatted screenplay.

TOOLS Phone voice notes / ChatGPT

02 Set up the project



One folder per project. Let the LLM build the subfolders; it reads whatever you drop in: PDFs, video, images.

TOOLS Claude / ChatGPT

03 Build the set in 3D



Give it a reference video, like a home tour. It pulls frames, models the space in Blender and checks its work against them.

TOOLS Claude + Blender connector

04 Block the scene



Feed it the script and the 3D layout. Get a top-down blocking diagram: entrances, paths, and where each character falls.

TOOLS Claude / ChatGPT

05 Cut an animatic



Have it animate rough characters in Blender with the script's action captioned underneath, like a studio animatic.

TOOLS Blender

06 Lock the look

Choose the style frame, palette, aspect ratio and camera and lens language before any shot is generated.

TOOLS ShotDeck / ChatGPT

GENERATION & POST

SHOTS / VIDEO / EDIT

07 Build the cast



Leads get separate high-res references per angle and expression. Extras come from one batch prompt. Layer on the natural-skin prompt.

TOOLS ChatGPT Images

08 Generate keyframes



Ask for all the shots that match the blocking, photoreal and in the locked style. Then give director's notes on what comes back.

TOOLS ChatGPT Images

09 Package the handoff



Numbered reference images plus a prompts.md file in one folder, so the next tool picks up without re-explaining.

TOOLS Claude > ChatGPT

10 Generate the video



Timecoded prompts, a reference legend, and 30-second clips joined with extend-forwards. Test at 480p first.

TOOLS Runway MCP / Seedance 2.5

11 Generate inserts



Make 30 seconds of 1–2-second cutaways in the same style to cover anything that still looks janky.

TOOLS Seedance 2.5

12 Assemble and cut

Let the AI stitch the clips into an assembly. Make the creative cut yourself.

TOOLS ChatGPT / DaVinci Resolve

WHERE THE WORKFLOW BROKE, AND THE FIX

Claude reached Runway through its API, which blocked shots that used Darion's own likeness. ChatGPT reached Runway through its MCP connector and went straight through. Class recommendation: run the whole workflow in ChatGPT for now.

LOOK DEVELOPMENT & PROMPTING

GET THE LOOK ON PURPOSE.

The habits, references and specs that make AI images read as cinema instead of stock.

PRINCIPLE 1

PROMPT LESS, THEN DIRECT


Today's models have absorbed Fincher, Tarantino and Deakins. Over-prompting crushes what they can offer. Start loose, review, then direct: "more close-ups," "keep the camera moving."

PRINCIPLE 2

ASK FOR WHAT YOU REALLY WANT

Don't shrink a request to the tool's known limits. Ask for the full one-minute scene and let the AI split it into clips, or show you where the wall is.

THE STYLE-REFERENCE RECIPE THE STEP MOST PEOPLE SKIP



An image beats adjectives: haze, palette and depth land closer when the model can see them.

TECHNIQUE SHEET ADD AS NEEDED

Color palette

Feed a palette from a real film frame. Left alone, ChatGPT drifts to a clean commercial look; cinema is graded.

Aspect ratio

Set it first. Darion prefers 2.40:1, a Deakins favorite. Seedance on Runway tops out at 21:9, the closest widescreen.

Coverage angles

Ask for the same shot from several angles, "each on a different angle of the 360° azimuth." The word does the work.

Negative space

AI fills every inch. Ask it to study the reference's composition and leave room for the eye to rest.

Camera & lens

Borrow real specs from ShotDeck metadata, e.g. Ex Machina: Sony CineAlta F65 with Cooke Xtal Express lenses.

Keep it running

ChatGPT generates several images at once and accepts new instructions mid-run. Keep adding notes as you think of them.

Depth

Ask for foreground, midground and background in every shot that allows it. Once it learns your taste, it keeps it.

Texture

Add "film grain, film noise, chromatic aberration" to pull images away from the polished AI look.

Ask how to ask

Stuck? Ask the model how to prompt it; that's how Darion found "azimuth." [Midjourney Describe](#) also turns images into prompt wording.

DEMO PROMPT · BUILT LIVE IN CLASS

"2.40:1 aspect ratio. A post-apocalyptic hunter standing over a zombie kill in a futuristic laboratory, a wash of red alarm light on his face. He's Jamaican and wearing leather. Use the cinematography specs and the style of the attached image. Give me that shot from seven different angles; make sure each image is on a different angle of the 360 azimuth. Add film grain, film noise and chromatic aberration to each shot."

Attached: an Ex Machina frame from ShotDeck (saturated red, night interior). No blocking given, on purpose: let the model propose first.

CHARACTERS & REFERENCES

BUILD THE CAST RIGHT ONCE.

Everything the video model makes is built on your references. Approve them before you generate a single shot.

REUSABLE PROMPT

The natural-skin layer

Identifiable pores

Baby hairs

Blemishes, even on models

Discoloration

Gradations of color

Darion's skin prompt has held up through nine months of model updates. Save it once, tell the LLM to fetch it from your chat history, then layer character specifics on top: sunken eyes, unruly hair.

BATCH, THEN FIX

Extras in one pass

"Give me eight thugs, some version of gothic. Throw some tattoos in there. Make sure they're diverse genders and races. Make them look grimy and greasy."

Generate all eight, then fix only the ones that miss. Save the detail work for your leads.

LEADS: SEPARATE REFERENCES BEAT ONE SHEET

SEEDANCE 2.5 TAKES UP TO 30 IMAGES

01

Full body
In costume

02

Front
Neutral, highest resolution

03

Three-quarter
Face

04

Side profile
Separate image

05

Back of head
Hair and shape

06

Expressions
Surprised, determined...

Character sheets were a workaround for 8–10-image reference limits. With 30 slots, upload each view as its own high-resolution image. More pixels per view means a sharper likeness, and a range of expressions lets a character duck, react and emote instead of wearing one look.

How to load them: put them at the top of the prompt and let the model work out who is who; your reference legend removes any doubt. Simple sheets are still fine for minor characters.

ONE WEAK REFERENCE, DAYS OF REWORK

Seedance copies what you feed it. A flawed hero sheet made Darion's face read heavy in every frontal shot: two to three days of regeneration. AI-generated Maserati references gave shots a cartoony cast that bled into others. **Use real photos for real products**, such as the manufacturer's site, and approve every foundation before you generate.

NUMBER EVERYTHING

Reference mapping

IMG 01 Style & cinematography reference

IMG 02 Darion (lead)

IMG 05 Buzz-cut man, table crew

IMG 08 Mohawk woman

Name files 01, 02, 03... and paste a legend like this into the video prompt, plus an identity and state lock: who is in the scene and who is already down.

PROFESSIONAL PROJECTS

Copyright hygiene

Distinctive buildings and designs can be protected. For paid or festival work, have the LLM alter a referenced design until it isn't a copy; it understands the request.


Only use images you have rights to. Firms now scan the web for unlicensed use.

General guidance from the session, not legal advice.

GENERATION & POST

MAKE IT MOVE, THEN MAKE IT CUT.

How the one-minute fight was generated in Seedance 2.5 through Runway, rescued with inserts and assembled.



Limits as used in class. Audio references: up to 10 advertised, one worked reliably for Darion. Check current limits before you plan.

SIX RULES FOR GENERATING FROM THE MODERN HOUSE BUILD

- Keyframes beat a blocking video**
The blocking animation didn't carry the action across. Photoreal keyframes did: the model hits each one and invents the motion between.
- Let ChatGPT write the timing**
It builds the Seedance prompt from your script and references, beat by beat: 0–2 s, 2–6 s, 6–8 s, on to 30 s.
- Extend forwards for length**
Generate clip 1 (0–30 s), upload it back as Video 1, and prompt an extension of exactly 30 new seconds.
- Spend like a producer**
Test at 480p; 1080p ran about \$13 per 30 seconds. Regenerate only what's broken: the crash-in and drive-out stayed, the middle minute was redone.
- Control the soundtrack**
Turn on "include audio" for effects and ask for no music. It doesn't always listen.
- Swap models mid-project**
When Seedance 2.5 shipped mid-way through a client job, switching delivered better results. Don't marry a version.

INSERTS COVER THE LAST 20%

AI gets you most of the way, and the remaining cuts nag. Crews shoot B-roll for exactly this: when a stray coffee cup lands in a hero shot, you cut to a detail and back. Generate the same safety net, then check every insert for AI errors (hands are the usual tell).

Edit like a studio. An assistant editor organizes footage and builds the assembly; the editor shapes it. ChatGPT can already stitch clips into that assembly. Make the creative cut in [DaVinci Resolve](#), and drop cool shots that don't make sense.

INSERT PROMPT · AS TYPED IN CLASS

"Some shots aren't working or moving quickly enough. To solve this, create a prompt for a 30-second Seedance 2.5 video of a series of relevant close-up and insert shots that I can use to intercut with the final video in DaVinci Resolve to hide anything that is still janky. Guns being cocked. Grabbing of a weapon. Thug-woman inserts on her weapon as she rounds the kitchen island. Sinister facial expressions of some of the thugs. Face-on shot of a gun pointed directly at camera, then firing. Rack focus from the muzzle of the gun to the face. Each shot need not be more than one to two seconds. Use the same style, cinematography reference and character references as needed."

WORKING METHOD

DIRECT IT LIKE A STUDIO.

Habits that keep a long AI project on track: memory, direction, judgment and budget.

SKILLS FILES

Project memory in .md files

Markdown files are plain-text notes any LLM can read. Keep the screenplay, style references, shot list, reference map and prompts in the project folder. New session? "Before we start, review the project folder." Long threads wear out; files don't.

WORK LIKE A DIRECTOR

Direct, don't operate

Don't prompt every image. Hand over the script, blocking, characters and style, ask for all the shots, then give notes the way a director would: "he seems too nonchalant," "keep these two exactly as they are."

REUSE WHAT WORKS

Turn lessons into skills

Have the LLM study strong examples, such as posters with great negative space, write what it learned into a skill file, then apply that skill to new work: shots, flyers, decks.

MODELS HAVE BIASES

Watch the defaults

In Darion's test, 995 of 1,000 Midjourney images for "attractive woman" were young white women with straight hair. Spell out age, ethnicity, hair and features, and test more than one model.

DON'T SKIP THE STEP

Storyboards: go straight to photoreal

Indies used to skip storyboards to save money. Keep the step, skip the sketches: a photoreal frame renders as fast as a sketch. Get soft approval, then generate video on day one.

TOOLS CHANGE IN WEEKS

Learn it, then break it

A workflow from three weeks ago is already dated. Use the one you're taught, then try to break it and find a better way. If a likeness drifts in a long thread, re-attach the reference.

BUDGET & CREDITS

SPEND WHERE IT SHOWS

Don't upgrade early

Stay on your ChatGPT plan until usage caps force a change; a few early resets come first. Darion used three in one week.

Trade posts for credits

Creator partnerships with [OpenArt](#) and [ImagineArt](#) swap posts for credits. Tools you reach through ChatGPT connectors can count.

The free route

[ComfyUI](#) is open source and runs locally, and ChatGPT or Claude can drive it. More setup, less spend.

HOMEWORK · BEFORE SESSION 2

- 1 Choose your capstone project for the four-week course.
- 2 Create its project folder for assets and .md files.
- 3 Generate images in ChatGPT to practice the workflow.
- 4 Hybrid filmmakers: research Beeble for live-action work.

CONTEXT & TOOLS

WATCH LIST & LINKS.

Where the industry is heading, and every tool mentioned in the session. Links checked October 5, 2026.

WHERE THE INDUSTRY IS GOING

CONTEXT FROM THE SESSION

Live-action / AI hybrid	Darion's prediction: most films become live-action/AI hybrids, the way nearly every production already uses some VFX, even if it's just wire removal.
AI gray box	Bitcoin (dir. Doug Liman, \$70M) was shot in 20 days inside a gray-screen box; AI supplies the sets, lighting and backgrounds. Producers say a practical shoot would have cost far more.
Tools to watch	Beeble relights and composites live-action plates for hybrid work. Intangible builds AI 3D sets for previs, though Blender driven by an LLM now covers much of that ground.
Creators to study	Gossip Goblin (Zack London): his first feature, <i>Gods Don't Give Gifts</i> , opens in U.S. theaters on October 30, 2026.
Vertical dramas	Phone-first micro-drama series (e.g., ReelShort) are booming, and social platforms are opening the format to creators.

TOOLS & LINKS

EVERYTHING MENTIONED, BY STAGE

WRITE & PLAN

ChatGPT [chatgpt.com](#)
The class pick for the whole workflow.

Claude [claude.ai](#)
3D builds, packaging and project organizing.

Final Draft [finaldraft.com](#)
Industry screenwriting software; optional for shorts.

3D & PREVIS

Blender [blender.org](#)
Free 3D suite for sets, blocking and animatics.

Blender connector [claude.com](#)
Setup guide for letting Claude drive Blender.

Blender MCP server [blender.org](#)
Official server for any MCP assistant.

Intangible [intangible.ai](#)
AI 3D scene building for previs.

LOOK & REFERENCE

ShotDeck [shotdeck.com](#)
Film stills with palette, camera and lens data.

Midjourney Describe [midjourney.com](#)
Turns an image into prompt wording.

GENERATE

Runway [runway.com](#)
Aggregator that runs Seedance and more.

Runway for ChatGPT [runway.com/mcp](#)
The MCP connector used in class.

Seedance 2.5 [bytedance.com](#)
ByteDance video model; 30-second clips.

Nano Banana Pro [deepmind.google](#)
Google image model, also on aggregators.

HYBRID VFX

Beeble [beeble.ai](#)
AI relighting and VFX for live-action plates.

EDIT & OPEN SOURCE

DaVinci Resolve [blackmagicdesign.com](#)
Edit, color and sound; free version.

ComfyUI [comfy.org](#)
Open-source, local, node-based generation.

CREDITS & PLATFORMS

OpenArt [openart.ai](#)
Multi-model platform with creator partnerships.

ImagineArt [imagine.art](#)
Image, video and voice suite.

Higgsfield [higgsfield.ai](#)
AI creative suite and model hub.

Magnific [magnific.com](#)
Creative AI suite, reachable via connectors.

WATCH & STUDY

Gossip Goblin [gossipgoblin.studio](#)
AI filmmaker; first feature Oct 30, 2026.

ReelShort [reelshort.com](#)
Vertical micro-drama app.

GLOSSARY

TERMS FROM THE SESSION.

Plain-language definitions, with a link where one helps.

Agentic

An AI that works through multi-step tasks on its own, like opening Runway and generating, while you add notes.

Aggregator
runway.com

One account for many AI models. Runway runs Seedance, Nano Banana Pro and its own models.

Animatic
studiobinder.com

Storyboard frames or rough animation cut to time, often captioned, to preview a scene's flow.

API vs. MCP
modelcontextprotocol.io

Two ways an AI reaches other tools: an API is a direct programmatic link; MCP is an open standard for assistants.

Aspect ratio

Frame width to height. 2.39–2.40:1 is cinema "scope"; AI video tools often cap at 21:9.

Assembly edit
studiobinder.com

The first cut: every scene in script order, minimal creative editing. Built by the assistant editor.

Azimuth
wikipedia.org

The horizontal angle around a subject, 0–360°. Naming it gets AI to rotate the camera.

B-roll
studiobinder.com

Supporting footage, such as details and cutaways, used to cover edits and hide mistakes.

Blocking
studiobinder.com

Where actors stand and move in a scene. A blocking diagram maps it from above.

Character sheet

One image showing a character from several angles. For leads, now replaced by separate high-res references.

Chromatic aberration
adobe.com

Color fringing along high-contrast edges from real lenses. Adding it makes AI images feel photographed.

Connector
help.openai.com

A ready-made link between an assistant and another service, such as Runway inside ChatGPT.

Extend forwards

Generating new seconds that continue an existing clip, using that clip as the reference video.

Four-walling
wikipedia.org

Renting theaters outright and keeping the box office: a form of self-distribution.

Gray box (AI)
nofilmschool.com

Filming live actors on a neutral gray-screen set, then building environments and lighting with AI.

Insert / cutaway
studiobinder.com

A short close-up of a detail, such as a gun being cocked, cut into a scene.

Keyframe

A locked image of a key moment. The video model invents the motion between keyframes.

Live-action / AI hybrid

Real performances and footage combined with AI-generated elements, the way VFX is used today.

LLM

Large language model, such as ChatGPT or Claude: reads your files, writes prompts, drives other tools.

Markdown (.md)
markdownguide.org

Plain-text files with light formatting that any LLM can read. Used here as project memory.

Micro-drama
televisionacademy.com

A phone-first vertical series of very short episodes built on cliffhangers.

Negative space

Empty areas of the frame that rest the eye and point it at the subject.

Rack focus
studiobinder.com

Shifting focus from one subject to another within a shot, such as muzzle to face.

Reference mapping

Numbering each reference image and listing who or what it is in the prompt, so the model never guesses.

Shooting script

The production draft where camera directions are added. The writer's script leaves them out.

Skill (file)

Saved instructions an LLM loads when needed, such as a style guide it wrote from examples.

Style reference
shotdeck.com

An image whose light, palette, haze and composition the model should match.

480p / 1080p

Video resolutions. Draft at 480p to save credits; render finals at 1080p.